

A360 Agent Exchange

Marketplace taxonomy, copy standards, bulk design production system,
and the end-to-end build plan

Core thesis: a curated, HIPAA-verified, EMR-neutral marketplace of named AI agents is the unclaimed distribution position in aesthetic medicine — and a realistic, fully-visualized store can be built in about three weeks for roughly \$200–300 in tools.

Prepared for	Aesthetics360 Agent Exchange build-out
Prepared date	July 2026
Contents	Taxonomy & schema, listing copy standards, asset production pipelines, 6-phase project plan
Use case	Execution reference for building out a360-v2-wse.vercel.app/exchange

This report is a planning and strategy document produced by AI-assisted research. Figures, vendor capabilities, and compliance postures require direct diligence before contracting or public use.

Contents

1. Categories (6)	5
2. Curated collections (homepage rows)	6
3. Badges	7
4. Tag vocabulary (use 4–7 per listing)	8
5. Pricing models (display patterns)	9
6. Listing schema (every listing in `/exchange/listings/` follows this)	10
7. Voice & style guide for all copy	10
8. Listing index (16)	11
1. The 10 rewrite rules (with before → after patterns)	12
Rule 1 — Lead with the buyer's number, not your technology	12
Rule 2 — Name the agent; give it a job title	12
Rule 3 — Taglines are ≤ 60 characters and contain tension	12
Rule 4 — Every feature bullet = verb + object + payoff	12
Rule 5 — Kill the adjectives; keep the mechanics	12
Rule 6 — Compliance is a selling feature, above the fold	12
Rule 7 — Price transparently, anchor to ROI, kill lock-in anxiety	12
Rule 8 — Show the KPI it moves, with a believable range	12
Rule 9 — FAQs answer the objection, not the softball	12
Rule 10 — One voice across the store	12
2. Structural upgrades for existing listings	13
3. Quick audit checklist (score each existing listing /10)	14
1. The core insight: screenshots are not an image-model job	15
2. Visual identity (one system, 16 accents)	16
3. Bulk production pipelines	17
Pipeline A — Screenshots (the realism backbone)	17
Pipeline B — Icons (1 hour total)	17
Pipeline C — Hero images (photoreal scenes)	17
Pipeline D — Videos (30s per agent, ~1 day for all 16)	17
Pipeline E — Marketplace chrome	17
4. Per-agent asset checklist (× 16)	18
5. Model/tool shortlist & budget	19

6. Realism details that sell the illusion 20

Phase 0 — Decisions & setup (Day 1) 21

Phase 1 — Data & copy into the app (Days 2–4) 22

Phase 2 — The mock-app template (Days 4–8) ★ critical path 23

Phase 3 — Brand assets (Days 6–9, parallel with Phase 2) 24

Phase 4 — Videos (Days 9–13) 25

Phase 5 — Store realism & polish (Days 13–15) 26

Phase 6 — Launch & follow-on (Week 4+) 27

RACI-lite (solo-founder version) 28

Risks 29

Definition of done 30

1. Categories (6)

#	Category	Slug	Icon concept	What lives here
1	Front Office	front-office	headset/phone wave	Reception, booking, speed-to-lead, webchat
2	Consult Room	consult-room	two chairs / speech bubble	Consult intelligence, scribe, simulation, financing
3	Revenue Growth	revenue-growth	upward spark line	Reactivation, nurture, membership, referral
4	Analytics & Attribution	analytics	radar/ledger	Attribution, yield, KPI benchmarking
5	Practice Operations	operations	clipboard/shield	Inventory, compliance, training
6	Marketing Studio	marketing-studio	megaphone/sparkle	Local visibility, content, social

2. Curated collections (homepage rows)

- **Staff Picks** — Sloane, Ledger, Ember
- **New on the Exchange** — most recent 4
- **The Revenue Recovery Stack** — Sloane + Ember + Fina (bundle story: "the second sale you already earned")
- **Runs Without Your EMR Changing** — everything with the `EMR-Neutral` badge
- **Built by A360** vs **Verified Partners** rows

3. Badges

Badge	Meaning	Visual
A360 Verified	Passed A360's PHI-flow review; BAA executed; no-training-on-PHI attested	teal shield w/ check
HIPAA BAA	Business Associate Agreement included at every tier	slate lock
Built by A360	First-party agent	gold monogram
Verified Partner	Third-party vendor, certified	silver handshake
EMR Sync	Live two-way sync (per-EMR chips: Boulevard, Zenoti, Aesthetic Record, PatientNow, Symplast)	plug icon + logo chips
EMR-Neutral	Works via CSV/API with any system	open plug
Outcome Pricing	Pay-per-result option available	tag with %
New / Staff Pick / Most Installed	Merchandising	pill labels

4. Tag vocabulary (use 4–7 per listing)

voice-ai, sms, webchat, booking, speed-to-lead, reactivation, recall, consult-intelligence, sales-coaching, attribution, roi-tracking, call-tracking, yield-management, scheduling, scribe, documentation, charting, simulation, visualization, financing, bnpl, memberships, loyalty, referrals, inventory, injectables, compliance, gfe, seo, local-marketing, reviews, content, social-media, before-after, multi-location, benchmarks, whatsapp, spanish, no-show-reduction, deposits, upsell, churn-prediction, kpi

5. Pricing models (display patterns)

1. **Flat monthly** — \$X/mo per location (most agents)
2. **Monthly + usage** — \$X/mo + \$Y/min or per message (voice-heavy agents)
3. **Per provider** — \$X/mo per provider (Quill scribe)
4. **Outcome** — \$X/mo base + \$Y per booked appointment (capped at \$Z/mo) — flagship differentiator; always shown with the audit note: *"Booked appointments are counted only when confirmed in your calendar — auditable in your Attribution ledger."*
5. Every listing shows: **14-day free pilot · Month-to-month · No setup fee** (deliberate anti-lock-in positioning) unless noted.

6. Listing schema (every listing in `/exchange/listings/` follows this)

- slug, persona name + role title, product name
- category, tags[], badges[]
- card: tagline (≤60 chars), short description (≤140 chars), price-from, rating (demo), installs (demo)
- hero: headline, subheadline, primary CTA, hero visual spec
- overview (150–250 words)
- key features (5–7, benefit-first)
- how it works (3–4 numbered steps)
- KPIs moved (2–4, with benchmark ranges – sourced from research where possible)
- integrations
- pricing tiers (2–3)
- proof bar (3 stat chips – demo data, marked)
- FAQs (4–6)
- sample reviews (2–3 – fictional, marked DEMO)
- compliance block
- media manifest: icon, hero, 4 screenshots, 30-sec video script

7. Voice & style guide for all copy

- **Named personas.** Every agent is an "AI employee" with a name, a job title, and a job description — the research showed named agents (Zenoti's Zeenie, Podium's Avery/Jerry, RepeatMD's Adonis & Aria) outsell feature lists.
- **Lead with the KPI, not the tech.** First sentence names the number that moves. Never open with "AI-powered."
- **Money is concrete.** Use practice-real figures: \$536 avg visit, ~\$285 CAC, 90-day tox cadence, \$4–12K/yr expired product.
- **Compliance is a feature, not a footnote.** Every listing states plainly: BAA included, PHI never used for model training, consent workflows built in.
- **Anti-lock-in is brand.** Month-to-month, EMR-neutral, "your data is always exportable, free" appears on every listing.
- Reading level: sharp operator, not clinician-formal. Contractions fine. No exclamation marks except in review quotes.
- **All ratings, install counts, reviews, and proof-bar stats in this package are demo/placeholder data** for the realistic-mockup goal; each is marked. Replace with real data before any public launch.

8. Listing index (16)

#	Persona	Product name	Category	Price from
1	Sloane	Consult Coach	Consult Room	\$499/mo
2	Ledger	Revenue Attribution	Analytics	\$299/mo
3	Ember	Patient Reactivation	Revenue Growth	\$199/mo + outcome
4	Tempo	Schedule Yield	Analytics	\$249/mo
5	Quill	Aesthetic Scribe	Consult Room	\$149/provider/mo
6	Remi	Front Desk	Front Office	\$349/mo
7	Piper	Lead Concierge	Front Office	\$249/mo
8	Fina	Financing Concierge	Consult Room	\$199/mo
9	Halo	Membership Optimizer	Revenue Growth	\$199/mo
10	Vista	Outcome Visualizer	Consult Room	\$299/mo
11	Beacon	Local Visibility	Marketing Studio	\$299/mo
12	Vial	Inventory Intelligence	Practice Operations	\$199/mo
13	Sentinel	Compliance Copilot	Practice Operations	\$249/mo
14	Muse	Content Studio	Marketing Studio	\$249/mo
15	Ripple	Referral Engine	Revenue Growth	\$149/mo
16	Lumen	Laser Parameter Copilot (LPOA)	Practice Operations → future Device Intelligence	\$249/mo + \$49/device pack

Copy Improvement Guide – Upgrading Existing Agent Exchange Listings

Note on access: the research environment's network policy blocks `vercel.app`, so the live listings at `a360-v2-wse.vercel.app/exchange`

1. The 10 rewrite rules (with before → after patterns)

Typical first-draft marketplace copy fails the same ways everywhere (Shopify's own listing guidelines, Salesforce AgentExchange top listings, and high-converting vertical-SaaS pages all converge on these fixes):

Rule 1 — Lead with the buyer's number, not your technology

- ✗ *"AI-powered patient engagement solution leveraging advanced LLMs"*
- ✓ *"40% of your leads call after 6pm. Remi picks up."*
- Test: does the first sentence contain a number the owner already worries about (\$285 CAC, 90-day tox cadence, \$536 visit)? If not, rewrite.

Rule 2 — Name the agent; give it a job title

- ✗ *"Automated Reactivation Module"*
- ✓ *"Ember — Patient Reactivation Specialist"*
- Named "AI employees" (Zenoti's Zeenie, Podium's Avery, RepeatMD's Adonis & Aria) consistently out-market feature names. Every listing should read like a hire, with a job description.

Rule 3 — Taglines are ≤ 60 characters and contain tension

- ✗ *"Comprehensive solution for practice growth and efficiency"* (says nothing)
- ✓ *"Stop paying for product that expires in the fridge."*
- Formula options: [pain, dry-stated] · [before vs after] · [the money left behind]

Rule 4 — Every feature bullet = verb + object + payoff

- ✗ *"Advanced analytics dashboard"*
- ✓ *"See margin per syringe, by provider — settle pricing arguments with a report, not a meeting"*
- If the bullet could appear on any product's page, it's not a bullet, it's filler.

Rule 5 — Kill the adjectives; keep the mechanics

- Strike on sight: *seamless, robust, cutting-edge, revolutionary, innovative, powerful, next-generation, holistic, effortless*. Replace with how it actually works: "answers from your synced catalog — if it's not on your menu, it doesn't improvise."
- Specificity is what makes AI products believable; vagueness is what makes them sound fake.

Rule 6 — Compliance is a selling feature, above the fold

- ✗ (HIPAA mention buried in FAQ #7 or absent)
- ✓ A visible compliance block on every listing: *"HIPAA BAA included · PHI never trains shared models · consent workflows built in."*
- In this vertical, the compliance paragraph closes deals that the feature list only opens. It's also the moat vs. GoHighLevel-wrapper competitors — say it loudly.

Rule 7 — Price transparently, anchor to ROI, kill lock-in anxiety

- ✗ *"Contact us for pricing"* (this is Recura/PatientNow's pattern — be the opposite)
- ✓ Table of 2–3 tiers + *"14-day free pilot · Month-to-month · Your data is always exportable, free."*
- Anchor line pattern: "One extra visit a month (\$536 average) covers the subscription twice."

Rule 8 — Show the KPI it moves, with a believable range

- ✗ *"Boost revenue!"*
- ✓ *"KPI moved: rebooking rate (analog benchmark: 15–25% of overdue patients per cycle)"*
- Ranges + named benchmark sources read as honest. Single big claims read as vendor math.

Rule 9 — FAQs answer the objection, not the softball

- Delete: "What is [Product]?" (that's the overview's job)
- Add the real ones: *Is recording patients legal? Will patients know it's AI? What happens when I cancel? How is a 'booked appointment' counted for billing? Who owns the data?*
- The billing-audit FAQ ("how do you count outcomes?") builds more trust than any testimonial.

Rule 10 — One voice across the store

- Sharp operator tone; contractions welcome; zero exclamation marks outside review quotes; sentence-case headings; the same section order on every listing (hero → overview → features → how it works → KPIs → pricing → FAQs → reviews → compliance). Consistency is what makes a marketplace feel real instead of assembled.

2. Structural upgrades for existing listings

1. **Adopt the shared schema** in `TAXONOMY.md` §6 — same sections, same order, every listing.
2. **Re-tag everything** against the controlled vocabulary (§4) — free-form tags kill filtering.
3. **Add badges** — especially `HIPAA BAA`, `EMR-Neutral / EMR Sync`, and `A360 Verified`; the badge row is the fastest trust signal on the card.
4. **Add the proof bar** (3 stat chips) to each listing — clearly marked demo data until real numbers exist.
5. **Cross-link agents** ("Works with: Sloane, Ledger") — the ecosystem story is a differentiator no single-product competitor can copy; every listing should sell one neighbor.
6. **Add 2–3 reviews per listing** (marked DEMO) — an empty review section reads as a dead marketplace; a marked-demo review section reads as a preview.
7. **Normalize pricing presentation** — "from \$X/mo" on the card; tier table on the page; pilot + month-to-month everywhere.

3. Quick audit checklist (score each existing listing /10)

#	Check	Pass looks like
1	First sentence has a buyer-relevant number	"\$285 to win a patient..."
2	Agent has a name + job title	"Ember — Reactivation Specialist"
3	Tagline ≤60 chars with tension	—
4	Zero banned adjectives	—
5	Features are verb+object+payoff	—
6	Compliance block above the fold	—
7	Transparent tiered pricing + pilot	—
8	KPI named with benchmark range	—
9	Objection-first FAQs incl. billing audit	—
10	Cross-links to ≥1 other agent	—

Anything scoring ≤6: rewrite from the templates in [exchange/listings/](#).

Design & Asset Production Guide — A360 Agent Exchange

Goal: produce a *realistic, live-feeling* marketplace — icons, hero art, product screenshots, and short videos for all 16 listings — fast, in bulk, and visually consistent. This guide gives the system, the tool choices, and copy-paste prompts.

1. The core insight: screenshots are not an image-model job

The single biggest realism mistake is generating "product screenshots" with Midjourney/DALL·E-class models — they produce blurry pseudo-UI with garbled text that instantly reads as fake. The professional approach:

Build real (mock) UIs in HTML and screenshot them.

You already have the perfect tool for this: Claude Code (or v0.dev/Lovable) generates a themed dashboard mock per agent in minutes, and Playwright screenshots them at pixel-perfect 2x resolution — identical fonts, identical chrome, real text, in bulk. This also gives you **free video**: screen-record the same mocks with cursor movement for launch clips.

Asset strategy by type:

Asset	Tool class	Recommended
Product screenshots (4–5 per agent)	HTML mocks + Playwright screenshot	Claude Code → one templated mock app, 16 themes
App icons (16)	Vector-native image model	Recraft (SVG output) or Ideogram; fallback: single prompt template in GPT-Image/Midjourney with fixed style block
Hero/lifestyle imagery	Photoreal image model	GPT-Image or Midjourney v7 (people/clinic scenes); Flux for speed
Marketplace UI chrome (cards, badges)	Design in code	Already exists in your exchange app — extend it
30-sec videos	Screen capture of mocks + AI b-roll + AI voiceover	Screen Studio (or Playwright video) + Veo/Runway/Kling b-roll + ElevenLabs VO
Diagrams ("how it works")	Code-rendered	Mermaid/HTML → screenshot, or Figma

2. Visual identity (one system, 16 accents)

Keep ONE design system so the store feels like one company; vary only an accent color and icon per agent.

- **Base UI:** dark-neutral dashboards (slate-900 background, white cards on slate-50 for light-mode shots), Inter or Geist type, 12px radius cards, generous whitespace. Matches modern medical-SaaS expectations (Boulevard/Mangomint aesthetic tier).
- **Brand base:** A360 deep teal `#0F766E` + ink `#0F172A` + ivory `#F8FAFC`. (For **Lumen**, reuse the LPOA doc palette — it already specifies `#0f766e` / `#2563eb` accents, amber `#f59e0b` review states, green `#16a34a` available, red `#dc2626` risk — designed for locked/review/available control states.)
- **Per-agent accents:** Sloane teal `#0F766E` · Ledger navy+green `#14532D` · Ember orange `#EA580C` · Tempo cobalt `#1D4ED8` · Quill ink-blue `#1E3A8A` · Remi coral `#F43F5E` · Piper violet `#7C3AED` · Fina emerald `#059669` · Halo gold `#B45309` · Vista lilac `#A855F7` · Beacon amber-gold `#D97706` · Vial glass-teal `#0D9488` · Sentinel graphite+amber `#475569` · Muse rose-gold `#BE7B6B` · Ripple aqua `#06B6D4` · Lumen spectrum-on-slate `#2563EB`.
- **Persona avatars:** do NOT use human faces for the agents (uncanny + legal simplicity). Use abstract "energy mark" avatars — the icon glyph in a soft-gradient circle. Distinctive, brandable, batch-generatable.

3. Bulk production pipelines

Pipeline A — Screenshots (the realism backbone)

1. **One mock-app template** (Next.js + Tailwind, same stack as LPOA Viewer doc): a dashboard shell with sidebar, header, and 4 swappable content panes driven by a JSON theme file (`agent`, `accent`, `panels[]`).
2. **16 theme files** — each agent's 4 screenshot panes are specified in its listing's *Media manifest* (e.g., Sloane: score card, missed-revenue ledger, provider leaderboard, mobile digest).
3. **Playwright batch script**: `loop themes → render → page.screenshot()` at 2880×1800 and 1170×2532 (mobile). ~64 pixel-perfect screenshots in one run.
4. Populate with the demo data already written in the listings (names, dollar figures, percentages) — consistency between copy and screenshots is what makes the store feel live.
 - Claude Code can build this template + themes + script in one session; it's the highest-leverage build task in the plan.

Pipeline B — Icons (1 hour total)

Master prompt (Recraft/Ideogram/GPT-Image), swap the [GLYPH]/[COLOR] per agent from \$2 and each listing's media manifest:

"Minimal flat app icon for a premium medical software marketplace: [GLYPH] as a single bold geometric mark, [ACCENT COLOR] on [BG], subtle 8% noise gradient, rounded-square container with 22% corner radius, no text, no gloss, no drop shadow, consistent 12% padding, vector style, 1024×1024." Batch all 16 in one session; regenerate outliers until the set reads as a family.

Pipeline C — Hero images (photoreal scenes)

Master prompt template (GPT-Image / Midjourney v7 `--ar 16:9 --style raw`):

"Editorial photograph inside a high-end medical spa, [SCENE from listing media manifest], soft natural window light, shallow depth of field, calm neutral palette with [ACCENT] accents, modern minimal interior, photorealistic, no text, no logos. Overlay-ready negative space on the [left/right] third." Then composite the matching mock-UI screenshot (Pipeline A) into the negative space as a floating panel — 20 minutes per hero in Figma, or automate with an HTML compositor template. **Never let the image model draw the UI.**

- People in shots: use clearly adult, diverse, consented-model-style imagery; avoid before/after realism for fictional patients (brand-safety) — imply outcomes, don't depict them.

Pipeline D — Videos (30s per agent, ~1 day for all 16)

Formula per the video scripts already written in each listing:

1. **[0–5s] Hook** — b-roll: Veo 3 / Runway Gen-4 / Kling from prompt: *"cinematic interior shot, upscale medical spa at [time], [hook scene], soft focus, premium commercial look, 4s"* (or stock from Artgrid/Pexels for zero-AI-artifact safety).
2. **[5–22s] Product** — screen-record the Pipeline-A mock with scripted cursor (Screen Studio for auto-zoom polish, or Playwright's video capture for full automation).
3. **[22–30s] Card** — end-card template in the accent color: persona name, tagline, "On the A360 Exchange."
4. **VO**: ElevenLabs, one consistent voice across all 16 (warm, unhurried, female-lean tested well for this vertical); scripts are in each listing.
5. **Assembly**: CapCut/Descript template — drop in the 3 segments + VO; export 16:9 and 9:16.

Pipeline E — Marketplace chrome

Cards, badge chips, category headers, collection rows: extend the existing exchange app's components. The badge system (§ TAXONOMY 3) needs 9 small SVGs — Recraft batch or hand-drawn in Figma in an hour.

4. Per-agent asset checklist (× 16)

- ☐ Icon 1024px (Pipeline B)
- ☐ Hero 2400×1350 composited (Pipeline C + A)
- ☐ 4 product screenshots @2x desktop + 1 mobile (Pipeline A)
- ☐ 30s video, 16:9 + 9:16 (Pipeline D)
- ☐ Proof-bar stat chips (text, already in listings)
- ☐ Avatar mark (icon glyph in gradient circle — derived from icon)

Full store extras: category header art (6), collection banners (4), empty states (2), OG/social share card template (1).

5. Model/tool shortlist & budget

Need	First choice	Alt	Est. cost (all 16)
Mock UIs	Claude Code + Next/Tailwind	v0.dev, Lovable	dev time only (~2–3 days)
Screenshots	Playwright	Puppeteer	\$0
Icons/badges	Recraft (SVG)	Ideogram 3, GPT-Image	<\$30
Hero photos	GPT-Image / Midjourney v7	Flux 1.1 Pro	<\$60
Video b-roll	Veo 3 / Runway Gen-4	Kling 2, stock (Pexels)	\$50–150
Screen recording	Screen Studio	Playwright video, Loom	\$89 one-time
Voiceover	ElevenLabs	PlayHT	<\$25
Assembly	CapCut template	Descript	<\$30

Total incremental cost for a fully-visualized 16-listing store: roughly \$200–300 + 4–6 focused days. The mock-app template (Pipeline A) is the critical path; everything else parallelizes.

6. Realism details that sell the illusion

1. Same fictional practice everywhere: "**Lumière Aesthetics — Austin, TX**", same staff names (Dr. Patel, Sarah RN, Maria at front desk), same patients (anonymized initials), same dollar figures as the listing copy.
2. Timestamps: recent weekday mornings; calendars showing realistic 70–90% booked density, not empty grids.
3. Numbers obey each other across screenshots (Ledger's totals = sum of Ember's + Remi's claims) — cross-checked in the theme JSONs.
4. Browser chrome in screenshots: none (app-frame only) — cleaner and avoids URL questions.
5. Every demo stat labeled subtly in-page footer: "Demo environment" — honest AND it reads like a real product's sandbox, which adds credibility.
6. A `DEMO DATA` watermark toggle in the mock template so you can flip between "presentation clean" and "public honest" exports.

Project Plan — A360 Agent Exchange: Full Realistic Build-Out

Objective: take the existing exchange at `a360-v2-wse.vercel.app/exchange` from prototype to a fully-realized, live-feeling marketplace: 16 complete listings with copy, taxonomy, icons, heroes, screenshots, and videos — demo-ready for prospects, partners, and investors. **Timeline:** ~3 working weeks solo; ~2 weeks with one designer/editor assisting. Costs: ~\$200–300 in AI/tool spend (see DESIGN_GUIDE §5). **Everything referenced below already exists in this repo** except the mock-app code and the media files themselves.

Phase 0 — Decisions & setup (Day 1)

#	Task	Output	Notes
0.1	Confirm store name & check trademark	Name decision	"Agent Exchange" collides with Salesforce's AgentExchange (their flagship brand since 2026 rebrand — see <code>research/MARKET_OPPORTUNITIES_2026.md</code>). For an internal demo it's fine; before public marketing, run clearance. Candidates: A360 Exchange , The Aesthetic Agent Store
0.2	Freeze taxonomy	<code>exchange/TAXONOMY.md</code> approved	Categories, badges, tags, collections
0.3	Freeze the 16-listing roster & demo data	<code>exchange/listings.json</code> approved	Adjust prices/personas to taste — one pass now, not during production
0.4	Pick tools & create accounts	Recraft/Ideogram, image model, ElevenLabs, Screen Studio (or Playwright video), CapCut	DESIGN_GUIDE §5
0.5	Export current live listings	<code>exchange/current-listings-export.md</code>	Enables the line-by-line copy rewrite (COPY_IMPROVEMENT_GUIDE) — the build environment can't reach <code>vercel.app</code>

Phase 1 — Data & copy into the app (Days 2–4)

#	Task	Output
1.1	Map <code>listings.json</code> fields → the exchange app's data model	Import script or seed file
1.2	Load all 16 card-level entries; wire category filters, badges, collections	Store renders 16 cards with search/filter
1.3	Build/extend the listing detail page to the schema (hero → overview → features → how-it-works → KPIs → pricing table → proof bar → FAQs → reviews → compliance block)	One template, 16 pages
1.4	Paste long-form copy from <code>exchange/listings/*.md</code>	All detail pages complete
1.5	Rewrite existing/legacy listings per <code>COPY_IMPROVEMENT_GUIDE.md</code> (10 rules + audit checklist)	Old items match new voice
1.6	Add "DEMO DATA" notice system (footer + toggle)	Honest realism

Exit criteria: the store is fully browsable with real copy and placeholder images.

Phase 2 — The mock-app template (Days 4–8) ★ critical path

#	Task	Output
2.1	Build ONE Next.js+Tailwind dashboard shell (sidebar, header, 4 swappable panes) driven by JSON themes	agent-mock/ app
2.2	Author 16 theme JSONs from each listing's Media Manifest (panel specs + demo data are already written per listing)	16 themes
2.3	Shared fictional world: "Lumière Aesthetics, Austin TX," staff, patients, cross-consistent numbers	world.json
2.4	Playwright batch: screenshot every theme (4 desktop @2x + 1 mobile each)	~80 screenshots
2.5	Lumen special case: the LPOA Viewer prototype (per LPOA_VIEWER_PROJECT_DOCUMENTATION.md) doubles as its mock — build it per that spec (PDF pane, chat, locked controls panel, confidence card) and screenshot the Scenario-A demo (Fitzpatrick IV / melasma / BBL)	LPOA Viewer V1 + 5 shots

Note: 2.5 is double-value — it produces the Lumen listing assets AND the standalone LPOA demo app the spec calls for.

Phase 3 — Brand assets (Days 6–9, parallel with Phase 2)

#	Task	Output
3.1	Batch-generate 16 icons + avatar marks (Pipeline B master prompt, per-agent glyph/color from DESIGN_GUIDE §2)	Icon set
3.2	Generate 9 badge SVGs + 6 category headers + 4 collection banners	Chrome set
3.3	Generate hero-scene photos (Pipeline C prompts per listing manifest)	16 hero plates
3.4	Composite screenshots into hero plates (Figma or HTML compositor)	16 finished heroes
3.5	OG/social share card template	1 template

Phase 4 — Videos (Days 9–13)

#	Task	Output
4.1	Generate/curate 16 hook b-roll clips (scripts in each listing's Media Manifest)	16 × 4–5s clips
4.2	Screen-record scripted walkthroughs of each mock theme	16 × ~15s product segments
4.3	Record VO from the 16 scripts (one ElevenLabs voice)	16 VO tracks
4.4	Assemble in one CapCut/Descript template; export 16:9 + 9:16	32 video files
4.5	Upload + embed on listing pages (player above screenshot carousel)	Live

Prioritization if time-boxed: videos for the 6 flagship/staff-pick agents first (Sloane, Ledger, Ember, Remi, Lumen, Vista); remaining 10 in a second batch.

Phase 5 — Store realism & polish (Days 13–15)

#	Task	Output
5.1	Reviews UI with the written DEMO reviews; rating stars on cards	Social proof layer
5.2	"Works with" cross-link chips between agents; bundle page for the Revenue Recovery Stack	Ecosystem feel
5.3	Install flow mock: "Add to my practice" → BAA step → EMR-connect step → success state	The killer demo moment: show <i>compliance-gated installation</i> , the thing no horizontal store has
5.4	Vendor pages (A360 + "Verified Partner" template)	Marketplace depth
5.5	Search, empty states, mobile pass, OG tags	Ship quality
5.6	QA pass: numbers consistent across copy/screenshots/videos; every DEMO marker in place	Sign-off

Phase 6 — Launch & follow-on (Week 4+)

- Demo script: a 10-minute guided tour (open store → Sloane listing → play video → install flow → Ledger proof ledger → LPOA live demo).
- Stakeholder/partner review sessions; capture feedback in this repo.
- Real-data path: as pilots go live, swap DEMO stats for actuals listing-by-listing (the schema doesn't change).
- Trademark clearance + public naming decision before external marketing.
- V2 candidates: Device Intelligence category buildout (more Lumen device packs), partner onboarding portal, checkout/billing.

RACI-lite (solo-founder version)

Workstream	Owner	Can delegate to
Copy & data (Ph 1)	You + Claude Code	— (done in this repo)
Mock app + screenshots (Ph 2)	Claude Code session	Any React dev
Icons/heroes (Ph 3)	You (1–2 days of prompt-batching)	Freelance designer (\$)
Videos (Ph 4)	You + template	Video editor (\$)
Store integration (Ph 1/5)	Whoever owns the a360-v2 app	Manus/dev partner

Risks

Risk	Mitigation
Existing app's data model doesn't match schema	<code>listings.json</code> is flat and renamable; map in one seed script (task 1.1)
Image models produce inconsistent icon families	Batch in one session with the fixed style block; regenerate outliers only
Demo data mistaken for real claims	DEMO markers everywhere + footer notice (tasks 1.6, 5.6)
Salesforce "AgentExchange" name conflict	Internal demo fine; clearance before public use (task 0.1)
Scope creep into a real billing marketplace	Not this project — install flow is a mock; billing is roadmap

Definition of done

A visitor can: browse 6 categories and 16 listings with consistent premium visuals → open any listing and see hero, video, 4 screenshots, pricing, FAQs, reviews → watch the install flow with BAA/EMR steps → and leave believing the A360 Agent Exchange is a live product one signup away. Every number they saw traces back to a file in this repo.