

A360 Market Opportunity Research

Agent Exchange & Product Portfolio — ranked opportunities, international white-label targets, and marketplace monetization models

Core thesis: no aesthetic-vertical platform runs an open, EMR-agnostic AI agent marketplace. Consultation intelligence and closed-loop revenue attribution are the two largest unclaimed white spaces, and both extend A360’s existing extraction engine.

Prepared for	Aesthetics360 leadership and partner discussions
Prepared date	July 2026
Method	Five parallel deep-research threads, ~130 web searches across ~160 sources
Use case	Agent Exchange product strategy, partner outreach, investor narrative

This report is a planning and strategy document produced by AI-assisted research. Figures, vendor capabilities, and compliance postures require direct diligence before contracting or public use.

Contents

Executive Summary	3
Part 1 — Ranked Product Opportunities	4
Tier 1 — Build (true white space + leverage on A360's existing extraction engine)	4
Tier 2 — Build on licensed rails (fast to ship, differentiate on integration + attribution)	4
Tier 3 — White-label / partner listings to fill the shelf	4
Explicitly avoid building	4
Part 2 — International White-Label / Exclusive-Distribution Candidates	7
Part 3 — Marketplace Model & Monetization	9
What the benchmarks say	9
Recommended model (sequenced)	9
Curation posture	9
Part 4 — Strategic Synthesis & Next Steps	10

Executive Summary

1. **The window is real but closing.** As of mid-2026 the incumbents (Zenoti's nine-agent "AI Workforce," Podium's Jan-2026 "AI Operating System for Aesthetics," Weave/TrueLark, Moxie's "Maia") are all shipping **closed, bundled AI agent suites** designed to deepen lock-in. **No aesthetic-vertical platform operates an open third-party app store.** An EMR-agnostic, curated agent marketplace has no direct incumbent — and first-mover in the marketplace also becomes the distribution channel for every product below.
 2. **The single biggest verified white space is consultation intelligence** — "Gong/Rilla for the consult room." Rilla hit ~\$70M ARR doing this for HVAC sales; Siro raised \$50M claiming +36% close rates and lists medspas as a *use case*, not a product. **No aesthetic-specific player exists**, HIPAA blocks the horizontal players, and A360 already owns the transcript-extraction pipeline. This is A360's most defensible product.
 3. **Second biggest: closed-loop revenue attribution.** No "Triple Whale for medspas" exists — ad spend → lead → consult → treatment revenue → LTV, closed against the practice ledger. HIPAA broke conventional pixels and B2B attribution tools won't sign BAAs. This matches A360's KPI obsession exactly and makes every other Exchange agent provable.
 4. **⚠ Naming flag: do not launch as "Agent Exchange."** In April 2026 Salesforce retired the 20-year AppExchange brand and unified its entire ecosystem (10,000+ apps, 1,000+ agents) under the name **AgentExchange**, backed by a \$50M builders fund. Run a trademark clearance and pick a defensible name (e.g., "A360 Agent Store," "A360 Exchange," "The Aesthetic Agent Market").
 5. **HIPAA is the moat, not a checkbox.** In nearly every category researched, the horizontal category leader is blocked or hobbled by PHI handling (recordings, pixels, patient lists, before/after imagery, review replies). The Exchange's core pitch: *"The AI tools proven in other industries, made legal for your practice — with the KPI lift measured for you."* An "A360 Verified" certification (BAA + SOC 2 Type II + no-training-on-PHI attestation + PHI-flow audit) is a differentiator no horizontal agent store offers — and, per the Epic model, the certification itself is monetizable.
 6. **Market context:** ~10,500–11,600 US medspas growing >\$1B/yr, ~\$1.4M average revenue per location, patient CAC ~\$285 and climbing, average visit ~\$536. Practices already paying \$300–\$1,000/mo for practice management will pay another \$300–\$1,500/mo for agents that demonstrably book revenue. Verified AI agent price points: receptionist \$200–\$600/mo, scribes \$99–\$600/provider/mo, "AI employee" add-ons \$99–\$399/mo.
-

Part 1 — Ranked Product Opportunities

Tier 1 — Build (true white space + leverage on A360's existing extraction engine)

1. Consult Coach — conversation intelligence for the consult room ★ flagship

- **Pattern proven by:** Rilla (~\$70M ARR, home services "virtual ridealongs," ~20% close-rate lift, real-time coaching since Sept 2025), Siro (\$50M Series B, +36% close rates, 30% lower rep turnover), Avoca Coach (QA-scores every CSR call; Avoca: \$125M+ raised, \$1B valuation, Apr 2026).
- **Aesthetic equivalent:** verified none. Zenoti's AI Scribe stops at charting; Podium's Avery stops at lead comms.
- **What it is:** record consults → score against a proven-closer rubric (talk ratio, objection handling, price presentation, financing-offer rate, package mention) → quantify un-offered opportunities in dollars mapped to the practice's actual catalog (A360's existing V3.1 extraction) → weekly injector/coordinator coaching → provider benchmarking.
- **Why A360 wins:** literally the front half of the existing Extraction Agent pipeline. HIPAA (PHI recordings, BAAs, 11 two-party-consent states) blocks Rilla/Siro from following.
- **KPI moved:** consult-to-close rate (+20–36% analog benchmarks), average ticket, treatment-plan acceptance rate.

2. Attribution Agent — closed-loop marketing attribution

- **Pattern proven by:** Triple Whale (e-commerce), HockeyStack (\$20M, Jan 2025), Invoca/CallRail (call tracking; CallRail has a HIPAA plan + BAA and a white-label/agency program), Patient Prism (dental — claims 30% revenue recovery from lost calls).
- **Aesthetic equivalent:** fragments only (agency GA4+CallRail dashboards, CorralData BI). No multi-touch attribution closed against EMR/POS revenue.
- **What it is:** HIPAA-safe pixel + call tracking + CRM/EMR join answering "which channel produced Botox revenue this month," with per-provider and per-location CAC, ROAS, LTV:CAC and budget reallocation recommendations.
- **Build vs white-label:** white-label CallRail rails; attribution logic is A360-built on its data layer. **Build this early — it's the connective tissue that makes every other Exchange agent provable.**
- **KPI moved:** cost per acquired patient (industry ~\$285), marketing ROI, LTV:CAC.

3. Reactivation / Recall Agent — treatment-cadence-aware patient win-back

- **Pattern proven by:** dental recall AI (practices carry 800–2,000 dormant patients; 15–25% reactivation per campaign; \$8K–\$25K/mo recovered production; 10–20x ROI), Impel in auto (recovers 33% of lapsed service customers), Keepme in fitness (churn prediction up to 95% accuracy).

- **Aesthetic equivalent:** partial (RepeatMD drop-off prediction; generic SMS blasts). Nobody does cadence-aware recall — tox is a metronomic ~90-day product, filler 9–18 months, laser packages fixed series — with text → text → AI-voice escalation and per-campaign revenue attribution.
- **Build vs white-label:** build targeting logic on A360 data; white-label the voice/SMS rails (Retell \$0.07/min self-serve BAA, Vapi \$0.05/min, Synthflow \$2K/mo white-label tier, VoiceAIWrapper \$499/mo with BAA).
- **KPI moved:** rebooking/recall rate, % of tox patients on-cadence, revenue per dormant patient. **The most quantifiable ROI story for launch.**

4. Yield / Utilization Agent — hotel revenue management for the appointment book

- **Pattern proven by:** hotel RevPAR systems; Book4Time (Agilysys) launched AI revenue management with dynamic treatment pricing for resort spas Q1 2025; Boulevard's "Precision Scheduling" optimizes slot placement but not price.
- **Aesthetic equivalent:** static rule engines only (Zenoti Dynamic Pricing, Salonized). Injector time is the scarce inventory; nobody optimizes revenue per provider-hour.
- **What it is:** score every open slot, forecast empty injector/laser hours, auto-fill via targeted offers to segmented waitlists (brand-safe, no public discounting). Coin the metric: **RevPAH — revenue per available appointment-hour.**
- **KPI moved:** provider utilization %, RevPAH. Completely unclaimed; pure-play for A360's KPI identity.

5. Aesthetic Scribe + Opportunity Extraction — "the scribe that pays for itself"

- **Pattern proven by:** vertical scribes beating horizontal ones in veterinary (ScribbleVet acquired by Instinct Jan 2026; Talkatoo 14K+ users w/ Patterson distribution); human-medical scribes at \$99–\$600/provider/mo; reviewers explicitly note most scribes lack cosmetic-specific documentation.
- **Race condition:** Zenoti shipped a medspa AI Scribe (Apr 2026) — window closing. Win on the fusion: ambient capture → aesthetic-native chart (units, lot numbers, GFE compliance language, consents) → **same audio feeds the extraction engine** (upsell candidates, package matches, future interests). No competitor links scribe → found revenue.
- **KPI moved:** documentation time, plus extraction-driven revenue per consult.

Tier 2 — Build on licensed rails (fast to ship, differentiate on integration + attribution)

6. **Front-Desk Voice Agent** — table stakes for the store; ~40% of medspa leads arrive 6pm–8am; the current supply is GoHighLevel-wrapper agencies (\$97/mo wholesale resold at \$200–\$600/mo) with weak HIPAA hygiene. Build on Retell/Vapi with real EMR-write, catalog-accurate pricing quotes, contraindication screening, deposit capture, and call-outcome attribution (Avoca's "voice-to-data" playbook). Avoid competing as "another receptionist" — win on integration depth.
7. **Speed-to-Consult / High-Ticket Nurture Agent** — leads contacted in 5 min convert 21x better; healthcare average response is 2h05m. Structurally offers a documented white-label API (its real-estate CRM partner generated \$3.4M new revenue in-app; integration <3 days). Add Manychat-style Instagram-

DM automation (the aesthetic demand channel) and multi-month nurture for surgical/\$5K+ pipelines (rhinoplasty leads convert 6–18 months out).

8. **Financing Orchestration Agent** — don't build a lender (Cherry: 60K+ providers; PatientFi: ~\$70M raised, Allergan/Galderma partnerships — saturated). Build waterfall routing across Cherry/PatientFi/CareCredit to maximize approvals, plus detect sticker-shock signals in consult transcripts and auto-trigger pre-qual. "Financing offered %" becomes a coached behavior inside Consult Coach.
9. **Membership Optimizer** — don't clone RepeatMD's storefront. The unserved piece is AI-designed membership economics (Clerri's dental playbook as software: member patients show 76% more visits, 146% more procedures): plan design from actual visit/spend data, breakage modeling, churn prediction, member-vs-non-member LTV.

Tier 3 — White-label / partner listings to fill the shelf

10. **Simulation Agent** — white-label **EntityMed** (aesthetic-native, explicitly white-label-able AI injectable/surgery outcome simulator + lead capture; one medspa: 120 qualified leads in a weekend). Then add the Overjet dental play on top: visual, priced, financeable treatment plans with per-provider acceptance-rate tracking (dental AI drove case-acceptance increases of 100%+).
11. **Visibility Agent** — white-label **Search Atlas / OTTO** (autonomous SEO agent; white-label included in \$99–\$999/mo plans) + a HIPAA-safe review-response layer (generic tools routinely violate HIPAA in review replies — the compliance nuance is the product). Includes AEO: showing up in ChatGPT/Gemini treatment queries.
12. **Supply / Margin Agent** — Torch Dental pattern (\$17M growth round; claims up to 32% supply savings). Forecast injectable burn from the booked calendar + historical units-per-treatment (A360 has this from extraction), optimize across Allē/Aspire loyalty economics, flag margin per syringe. Average medspa loses \$4–12K/yr on expired product. Contested (Zenoti AI Inventory Manager, Moxie GPO) but no independent player.
13. **Compliance Copilot** — always-on agent auditing charts for missing GFES/co-signs, scope-of-practice, state rule changes; PE-due-diligence readiness. Qualiphy (\$27.99/exam) and Spakinect commoditized the exam itself; continuous monitoring is unowned (only Moxie gestures at it, only for its own customers).
14. **Content / Ambassador Agents** — consent-gated before/after clip pipeline with e-signed media releases attached to every asset (patient likeness = PHI; the consent workflow is the moat vs OpusClip-class tools); patient micro-ambassador programs with FTC + med-board compliance. Later-stage shelf fill.

Explicitly avoid building

Another membership storefront (RepeatMD/PRIVI), another BNPL lender (Cherry/PatientFi), another undifferentiated AI receptionist (10+ players, being commoditized by Zenoti/Podium/Weave bundling).

Part 2 — International White-Label / Exclusive-Distribution Candidates

Companies with proven home-market products, explicit API/SDK/white-label programs, and **no meaningful US aesthetic-clinic presence** (verified). Suggested outreach order:

#	Company	HQ	What	Why now
1	Skinive	Netherlands	CE-marked AI skin analysis (50+ conditions, 3M+ images) with an explicit white-label API	Small company = fast deal; position as cosmetic/wellness triage (no FDA claims)
2	Arbrea Labs	Zurich (ETH spinout)	iPad-native 3D/AR face + breast/body outcome simulation, no camera rig, ~90 seconds	Crisalix (its main rival) was just PE-acquired (Mar 2025) and is already US-established; Arbrea is partner-hungry — US-exclusive distribution ask
3	Healing Paper / Gangnam Unni + Syscode; also KOS	Seoul	Korea's #1 aesthetics platform (\$32M Series C; ~\$73M rev +45% YoY) + subsidiary Syscode's clinic CRM — the OS of the world's densest aesthetic market	Korean CRM playbooks (recall cadence, package upsell, membership economics) are exactly A360's revenue-intelligence IP; expansion focus is Japan/SEA, not US. Licensing/JV framing; Korean-language diligence needed
4	Haut.AI	Estonia	150+ facial biomarkers + SkinGPT photorealistic treatment-outcome simulation; no-code branded embed	"Show the patient their skin after treatment" = consult-conversion tool; sells to beauty brands, not US clinics
5	Lululab	Seoul (Samsung spinout)	LUMINI AI skin-analysis kiosks + SDK (5M+ data points); clinic push at KHF 2025	"K-beauty clinical AI" positioning + lobby kiosk feeding the A360 platform
6	Leona Health	Mexico City	WhatsApp-first AI co-pilot for doctors; \$14M seed led by a16z (late 2025) ; 14 countries	WhatsApp-native patient comms for Spanish-speaking US markets (Miami/TX/SoCal); move before a16z pushes them upmarket in the US
7	medicalforce	Tokyo	#1 aesthetic-clinic EMR/CRM in Japan (~400+ clinics, ~\$10M Series B); zero US presence	Self-pay billing + multi-session "course" selling logic; longer-horizon licensing/JV. (Corridor validated: Japan's SBC Medical invested in US medspa platform OrangeTwist, Jan 2026)

#	Company	HQ	What	Why now
8	Binah.ai / Cherry Imaging / Orbo AI	Israel / Israel / India	Video-based vitals SDK (white-label) / TRACE 3D outcome quantification / BeautyGPT white-label skin-AI + AR API (~\$4.4M raised, multi-ethnicity training data)	Differentiator add-ons: wellness intake for GLP-1 programs; quantified before/after proof; budget-priced skin AI

Already in the US (integration partners maybe, exclusivity off the table): Crisalix, FotoFinder, QuantifiCare, Perfect Corp, Pabau, Phorest, Fresha, Booksy, Fresh Clinics, Revieve. **Watch list:** SkinVision (FDA pending — potential first-mover US distribution deal post-clearance). Also: a UK/EU cluster of small GDPR-native aesthetic AI receptionists (Parnidia, Mustly, Aesthetic Telephone AI) as cheap acqui-license targets for trained vertical conversation models.

Regulatory notes for any deal: simulation/visualization and cosmetic-grade scoring are generally non-device (no FDA claims); lesion/disease detection is device territory (510(k)/De Novo or strip the claims). Every foreign partner needs a BAA + US data residency — **A360's compliant platform is a negotiating asset: partners get US compliance "for free" by going through you.** WhatsApp requires Business API via a HIPAA-capable BSP. Face scans trigger Illinois BIPA / TX / WA biometric statutes — push consent/retention into the integration layer.

Part 3 — Marketplace Model & Monetization

What the benchmarks say

- **Shopify:** 0% take under \$1M lifetime / 15% above — low take as supply subsidy; monetize the merchant.
- **Salesforce AppExchange → AgentExchange:** 15% of net revenue (25% OEM); \$999 security review — **certification as trust moat.**
- **ServiceTitan:** annual partner dues + per-connection fee OR rev share — the proven mid-market vertical design.
- **Toast:** up to 30% of partner revenue in perpetuity — **owned distribution justifies high take rates.**
- **Epic:** ~\$1,900/yr vendor membership + \$500/yr listing — in healthcare, **the vetting is the product.**
- **Agentforce pricing lesson:** pure per-conversation pricing (\$2/conv) confused buyers; hybrid base + metered (Flex Credits, \$0.10/action) won. **GPT Store lesson:** engagement-pool payouts fail for B2B; per-seat/per-outcome contracts work.
- Marketplaces that own billing rails capture 15–30%; listing-only directories capture almost nothing.
- Adopt **A2A/MCP** as the integration standard so third-party agents list with minimal custom work.

Recommended model (sequenced)

1. **Launch bundles first (anchor tenants):** A360-branded agent tiers with 50–60% wholesale rev-share underneath — e.g. Front Desk AI \$349/mo; Growth Suite \$749/mo (adds Consult Coach, Reactivation, Visibility); Enterprise \$1,500+/mo/location. At 500 Growth Suite practices ≈ \$4.5M ARR, ~\$2M gross profit. Fastest revenue, de-risks the store launch.
2. **Open the curated store:** 10–25 "A360 Verified" third-party agents. Vendor side: \$1,500–\$2,500/yr certification fee (BAA execution, PHI-flow security review, sandbox) + **25% of net revenue billed through A360** (justified vs AppExchange's 15% by owned distribution + HIPAA wrapper — Toast proved 25–30% clears). 500 practices × 1.8 agents × \$400/mo × 25% ≈ **\$1.08M/yr**; at 2,000 practices ≈ \$4.3M/yr.
3. **Add outcome pricing once attribution is instrumented:** \$99–\$199/mo base + \$8–\$12 per agent-booked appointment (capped ~\$600/mo for bill predictability). At \$536/visit that's a ~2% take on driven revenue — easy ROI story. Define "booked appointment" contractually and auditably (the Sierra/Decagon dispute lesson). Layer a **benchmarking/analytics upsell** (\$99–\$249/mo): cross-practice agent-performance benchmarks — near-100% margin, unique to whoever owns the rails.

Curation posture

Small, deeply vetted catalog beats a big open one in a trust-sensitive vertical. "A360 Verified" = BAA in place + SOC 2 Type II + no-training-on-PHI attestation + PHI-flow architecture review. Own billing; never let third parties own the practice relationship.

Part 4 — Strategic Synthesis & Next Steps

The thesis in one line: incumbents are converging on closed AI bundles for scheduling and communications, while the clinical-revenue core of the business — *what was said in the consult, what it's worth, who converts, what's expiring in the fridge, and whether the chart is compliant* — remains spreadsheet-ware. A360's extraction engine attacks exactly that core, and an open agent marketplace is the unclaimed distribution position.

Recommended 90-day actions:

1. Trademark clearance + rename the marketplace (Salesforce AgentExchange conflict).
2. Productize **Consult Coach** from the existing extraction pipeline (fastest path from current assets to flagship product).
3. Instrument **Attribution** (CallRail white-label + EMR join) — it makes every other agent's ROI provable and is the KPI story your distribution partner sells.
4. Ship **Reactivation Agent** on white-labeled voice rails (Retell/Synthflow) — the most quantifiable launch ROI (\$8–25K/mo dental analog).
5. Open white-label conversations: Skinive, Arbrea, EntityMed, Structurely, Search Atlas/OTTO (all have explicit programs); begin Korean diligence on Syscode/KOS.
6. Define the "A360 Verified" certification program and adopt A2A/MCP as the listing standard.
7. Boulevard's open GraphQL API and Aesthetic Record's connector posture are the natural first integration rails for the EMR-agnostic story.

Appendices (full sourced memos):

- `appendix/01-cross-industry-ai-agents.md` — dental, vet, auto, hospitality, legal, fitness, home services patterns + white-label supply chain
- `appendix/02-international-white-label-candidates.md` — EU/UK/Israel/Korea/Japan/LatAm/ANZ companies + regulatory notes
- `appendix/03-sales-marketing-tech-transfer.md` — 12 category deep-dives with KPI mapping
- `appendix/04-aesthetic-software-landscape.md` — competitor tiers, funding, complaints, ranked white-space gaps
- `appendix/05-marketplace-business-models.md` — take rates, certification economics, agent pricing models